




Nexus AI — Agentic AI agents for the enterprise

Build · orchestrate · govern AI agents in production

Nexus AI

The enterprise platform for AI agents you can trust.

Project Context

Nexus AI is an enterprise platform for building, orchestrating, and governing AI agents in production — across both customer service and employee work. Enterprises love the promise of AI agents, but stall on the gap between a flashy demo and a system that survives real production, audits and regulators. Nexus AI was designed to close that gap: turning agent development from a risky, months-long engineering effort into a fast, governed, repeatable process. The design challenge was to make production-grade concerns — testing, orchestration, audit, compliance — feel native and effortless, not bolted on. Every screen is built to earn the trust of the people who sign off on enterprise AI.

CLIENT	ROLE
Nexus AI (enterprise platform)	Product Design & Build
TEAM	TIMELINE
Individual Project	12 Weeks · 2026

Problem

Enterprises kept hitting the same wall: dozens of promising AI pilots, almost none in production. Teams stitched together five or six disconnected tools to ship a single agent — with no shared governance, no audit trail, and no way to prove to a regulator why the AI did what it did. Compliance blocked launches, engineering became the bottleneck, and AI budgets burned with little to show. In regulated industries — banking, healthcare, telecom — a wrong or unexplainable AI decision isn't just embarrassing; it's a legal and financial risk. **They needed to build, orchestrate and govern agents in production — fast, safely, and without betting the company on a black box.**

~80%

of enterprise AI pilots never make it into production.

5×

engineering overhead — a team of five to ship what one developer should.

months

to move a single agent from prototype to a governed production system.

Indicative figures. Sources: industry analyst reporting on enterprise AI adoption · practitioner interviews. Replace with your study's data.

Approach

I designed the platform around the people who actually approve enterprise AI — engineers, compliance, and the business. The methodology: treat production-readiness as the product, not an afterthought, and make trust visible at every step rather than asking anyone to take the AI on faith.

1

Built for production, not demos

Every workflow can be tested and validated before it ever reaches a customer, so what ships has already been proven to work.

Key decision: prove it works before you deploy it.

2

Governance as the default

Audit, compliance and guardrails are enforced at the runtime layer — 100% of decisions logged and explainable, not left to the model.

Key decision: compliance the platform enforces, not the AI interprets.

3

Orchestrate, don't rebuild

A visual canvas lets teams compose and run purpose-built agents in parallel — reusable, LLM-agnostic, no lock-in.


Key decision: one platform for scripted + reasoning AI.

What was the problem I was trying to solve?

How might we let enterprises build, orchestrate and govern AI agents in production — fast, safely, and without vendor lock-in or a leap of faith?

Final Solution

One platform to ship agents to production, orchestrate them at scale, and prove they're working — safely.



Agent Platform

Every agent in production, in one view — with **live status, interactions and automation rate** at a glance, and a real-time health-and-governance panel beside them.

Uptime, compliance flags and explainability are first-class, so the people who own the risk can see the system is sound without digging.

Agents in production

Live metrics

Governance health

Visual Agent Builder



A drag-and-drop canvas lets teams **compose and orchestrate multi-agent workflows** — a router, knowledge and action agents, tools and conditions — with no code.

Purpose-built agents run in parallel with independent fault recovery, so real business complexity is handled, not avoided.

No-code canvas

Multi-agent

Parallel orchestration



Analytics & Governance

Outcomes are tied to the work: **automation rate, interactions and cost saved** trend over time, so value is provable, not anecdotal.

A compliance panel makes audit-readiness obvious — 100% logged, full audit trail per decision, SOC 2 / GDPR / HIPAA ready.

Outcome analytics

Full audit trail

Compliance-ready

Results

The platform turned stalled pilots into governed agents in production — with the numbers to prove it.

10×

faster delivery — months to days

71%

of customer & employee interactions automated

100%

of decisions logged, audited & explainable

\$2.4M

estimated annual cost saved per deployment

Representative enterprise outcomes — replace with your project's verified metrics.

“

Nexus AI took us from one FAQ bot to forty governed agents in production in a single quarter — with audit trails our regulators actually accept. It removed the engineering bottleneck and let us scale AI without scaling risk.



VP, AI Transformation

Global bank · client

Tools & Timeline

Twelve weeks from first workshop to a governed, production-ready platform — designed and built end-to-end.



Weeks 1–2

Discovery

Stakeholder workshops, compliance constraints, success metrics.



Weeks 3–5

Design

Flows, design system, dark + light UI, orchestration UX in Figma.



Weeks 6–10

Build

Front-end build, agent canvas, AI integration, iterative testing.



Weeks 11–12

Pilot & Launch

Governed pilot, audit hardening, production release.

Tools & technologies

Figma

React

TypeScript

Next.js

Tailwind CSS

LLM / AI APIs

Vector DB

Supabase

Scope

Product strategy · UX flows & orchestration model

Design system · enterprise UI, dark + light

Front-end build · agent canvas & AI integration

Outcome & Reflection

A production-grade enterprise platform that turns AI agents from risky pilots into governed, scalable systems — compressing months of engineering into days, with trust and compliance built in from the first screen.

WHAT I'D EXPLORE NEXT

Deeper agent observability, a reusable agent marketplace, and outcome attribution — tying every point of EBIT gain back to a specific agent optimisation, so the platform proves its own ROI.